



# Working with Jupyter on the Clusters

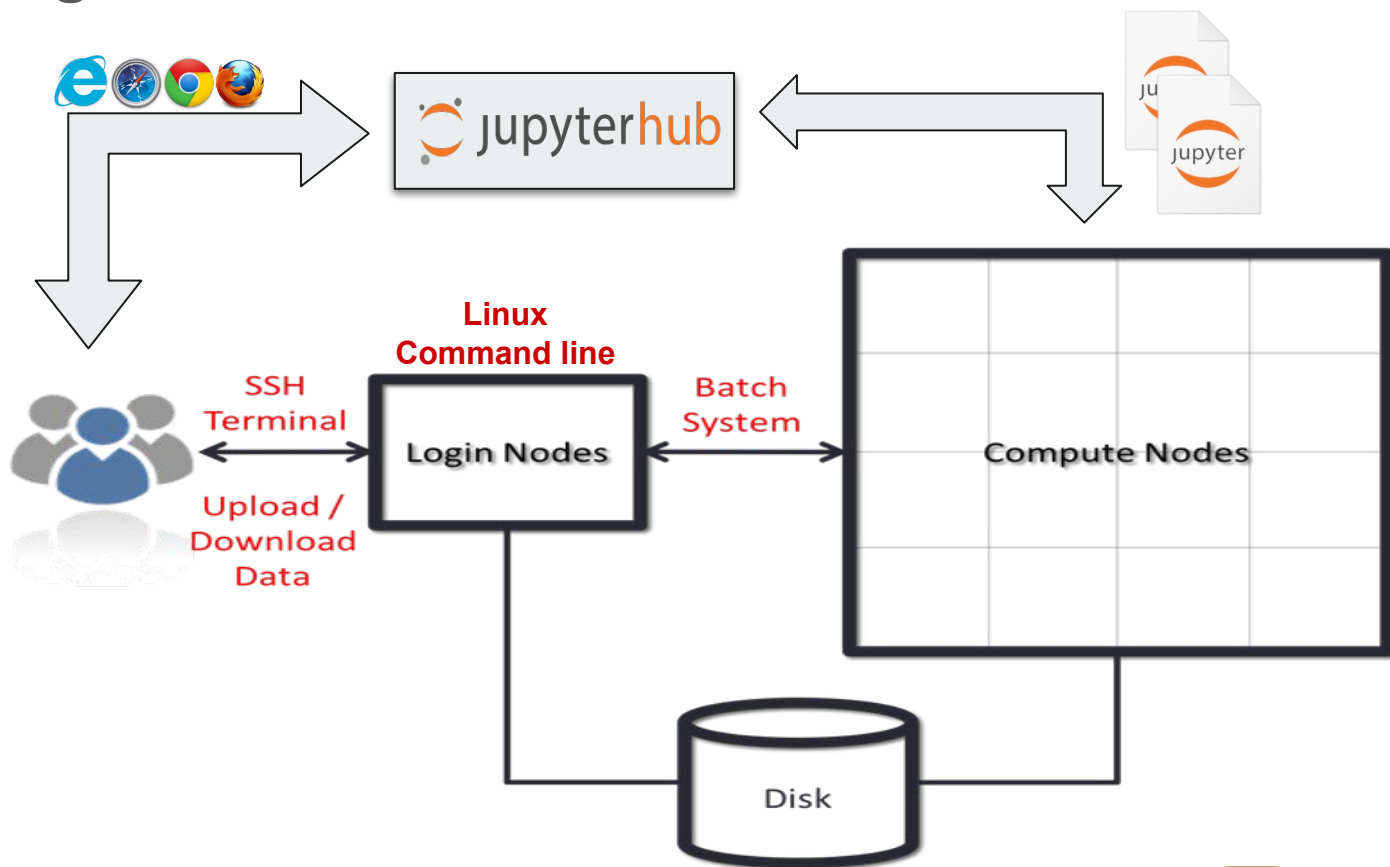
Jinhui Qin

SHARCNET | Compute Ontario | Digital Research Alliance of Canada  
Western University

# Outline

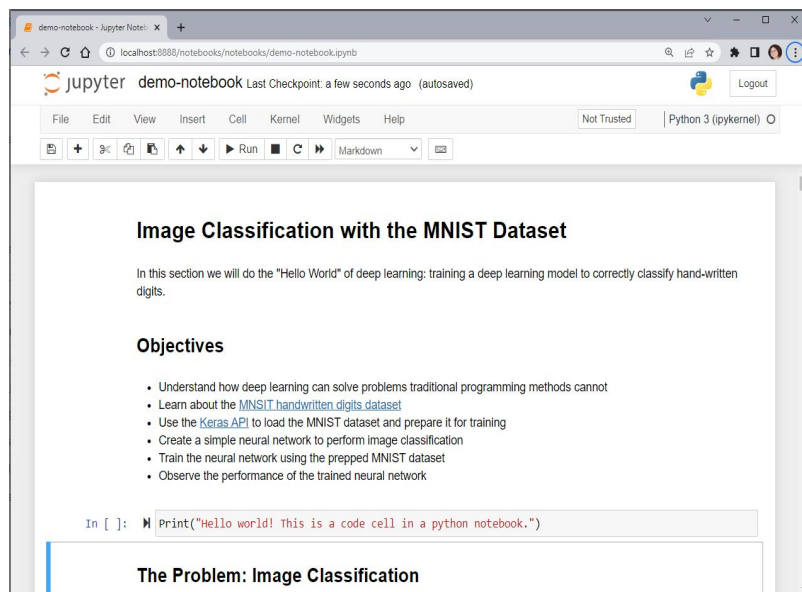
- Working on a cluster
- Notebook interface
  - Jupyter Notebook
  - JupyterLab
- Launching a Jupyter server
  - Via JupyterHub
  - Via Command line
- Use cases to demo
  - Case 1 : via JupyterHub without command line
  - Case 2: work with a commercial software package
  - Case 3: work with a container

# Working on a cluster

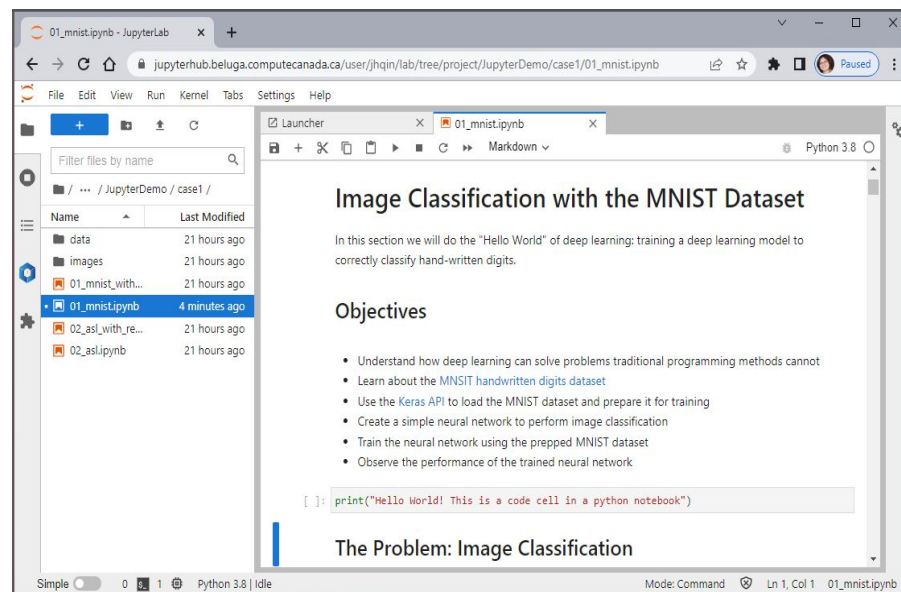


# Notebook Interface

## Jupyter Notebook — the classic interface



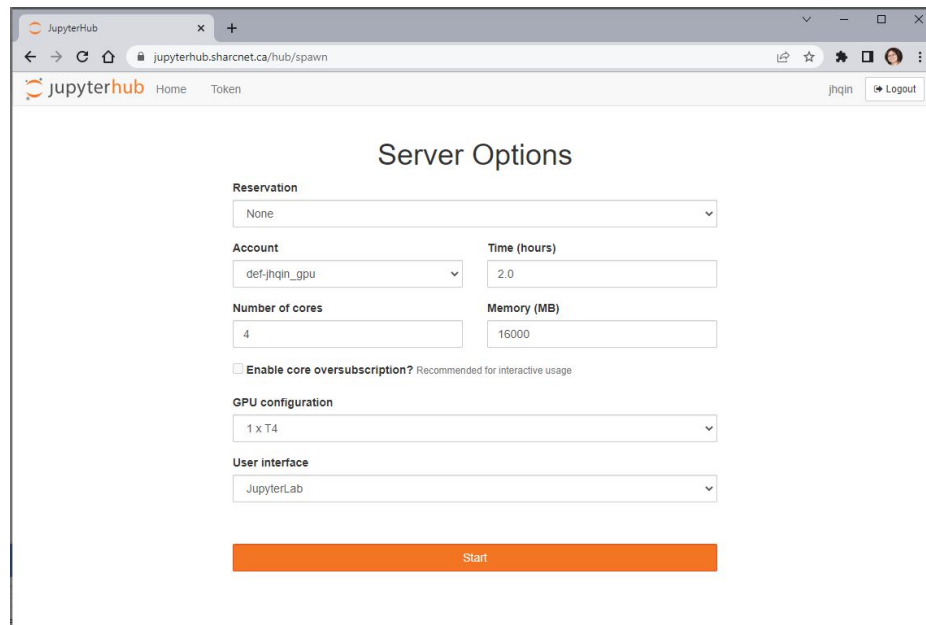
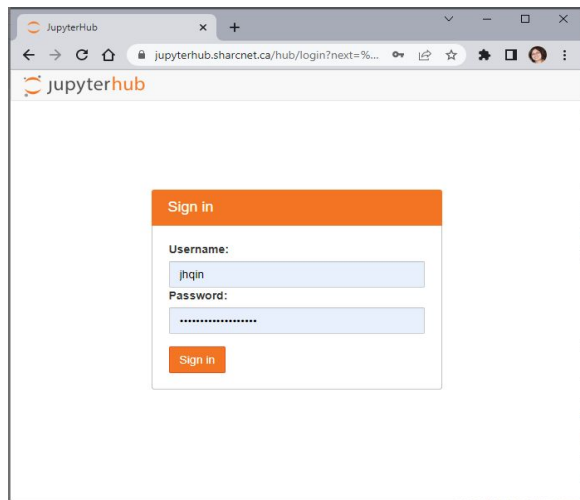
## JupyterLab — the new interface



# Launching a Jupyter server on a cluster

Via **JupyterHub** -- a web portal for launching a Jupyter server

- [https://docs.alliancecan.ca/wiki/JupyterHub#JupyterHub on clusters](https://docs.alliancecan.ca/wiki/JupyterHub#JupyterHub_on_clusters)



# Launching a Jupyter server on a cluster

## Via Command line

- ssh to a cluster
- load a **python** module and other modules if needed, e.g. scipy-stack
- create a python virtual env and Install **jupyterlab** and other libs if needed
- request an interactive session on a compute node
- launch a Jupyter server via command **jupyter-lab** from the virtual env.

```
[you@local]$ ssh user@cluster.alliancecan.ca
```

```
[user@cluster-login]$ module load python scipy-stack <other-modules>
[user@cluster-login]$ virtualenv --no-download ENV
[user@cluster-login]$ source ENV/bin/activate
(ENV)[user@cluster-login]$ pip install jupyterlab <other-libs>
(ENV)[user@cluster-login]$ deactivate
```

```
user@cluster-login]$ salloc --account=def-user --cpus-per-task=4 \
--mem=16G --gpus-per-node=t4:1 --time=1:0:0
[user@node###]$ module load python scipy-stack <other-modules>
[user@node###]$ source ENV/bin/activate
(ENV)[user@node###]$ jupyter-lab --no-browser --ip $(hostname -f)
```

# Features

	JupyterHub	Command line
Connecting	Automatically	<i>Need to setup a <a href="#">SSH tunnel</a></i>
Default environment	Pre-configured, <a href="#">temporary</a>	Self-made, <a href="#">persistent</a>
Internet access from a notebook	No	No
Accessing files on a cluster	Yes	Yes
Installing packages	Only those from system wheelhouse, i.e. with <i>--no-index</i>	Possible from external, e.g. <a href="#">PyPI</a> , on a login node
Working with a self-built virtual env	<a href="#">Possible</a>	Yes
Using system modules	Yes	Yes
Using GPUs	Yes	Yes
Working with a container	<a href="#">No</a>	<a href="#">Yes</a>



# Demo on Cluster Case 1







# Use Cases

## Case 1: run a python notebook for image classification on GPUs

- Launch a JupyterLab on a GPU node from a [JupyterHub](#)
- Work in the pre-configured environment
- Load modules, e.g. scipy-stack for using numpy, pandas, matplotlib, etc.
- Install additional python libraries, e.g. tensorflow
- Check GPU status if available
- Upload / Download / Export files
- No need to work in command line on a terminal

JupyterHub on clusters: [https://docs.alliancecan.ca/wiki/JupyterHub#JupyterHub\\_on\\_clusters](https://docs.alliancecan.ca/wiki/JupyterHub#JupyterHub_on_clusters)

# Notes for working via JupyterHub

- If any system module is needed, always load the system modules before installing any additional python packages, and kill all running kernels
  - Click  to open the module panel and load a system module
  - Click  to open the kernel panel and kill all running kernels
- Before working on a second notebook, always kill any previous running kernel
  - Click  to open the kernel panel and kill any previous running kernel
- After a python package installation, always restart the kernel before working on a notebook
  - Click  on top of the notebook panel to restart the kernel
- You can only install a python package with **--no-index** option on a notebook, e.g.

```
[ ]: !pip install --no-index scikit-learn
```





# Demo on Cluster Case 2

# Use Cases

**Case 2:** run a notebook to work with a commercial package via a *persistent* python virtual environment

- [An example](#) notebook to work with [Gurobi](#) optimizer with machine learning
- Prepare a python env with required packages installed
  - Required python packages:
    - pandas, numpy, matplotlib i.e. `module load scipy-stack`
    - seaborn, scikit-learn i.e. `pip install --no-index seaborn scikit-learn`
    - gurobipy, i.e. follow the instructions on [docs wiki](#)
- Two options for launching a Jupyter Server, ie.
  - via JupyterHub or via Command line

# Use Cases

**Case 2:** run a notebook to work with a commercial package via a *persistent* python virtual environment

- **Option 1.** launching a Jupyter server from a [JupyterHub](#)
  - Add a customized kernel from *command line on the cluster*

```
(gurobiENV)[user@cluster-login]$ pip install --no-index ipykernel
(gurobiENV)[user@cluster-login]$ python -m ipykernel install --user --name gurobiEnv
```

- Check available Jupyter kernels, remove a kernel if no longer needed

```
(gurobiENV)[user@cluster-login]$ jupyter kernelspec list
(gurobiENV)[user@cluster-login]$ jupyter kernelspec uninstall <kernel-name>
```

# Use Cases

Case 2: running a notebook for using a commercial package via a python virtual environment

- **Option 2.** launching a Jupyter server from **command line** on a compute node
  - Need to install *jupyterlab* in the same virtual env

```
(gurobiENV)[user@cluster-login]$ pip install --no-index jupyterlab
(gurobiENV)[user@cluster-login]$ deactivate

[user@cluster-login]$ salloc --account=def-user --cpus-per-task=4 --mem=16G --time=1:0:0

[user@node###]$ module load python scipy-stack gurobi
[user@node###]$ source gurobiENV/bin/activate
(gurobiENV)[user@node####]$ jupyter-lab --no-browser --ip $(hostname -f)
```

# Use Cases

Case 2: running a notebook for using a commercial package via a python virtual environment

- **Option 2.** launching a Jupyter server via command `jupyter-lab` on a compute node

```
(gurobiENV)[user@node###]$ jupyter-lab --no-browser --ip $(hostname -f)
... ..
http://<node-hostname>:8888/lab?token=109c369980ab346bcd129eb...
```

- Setup a SSH tunnel on a second terminal from your local computer

```
[you@local]$ ssh -L 9999:<node-hostname>:8888 user@cluster.computeCanada.ca
[user@cluster-login]$
```

- The URL to connect on a local web browser would be:

<http://localhost:9999/lab?token=109c369980ab346bcd129eb...>



# Demo on Cluster Case 3



# Use Cases

## Case 3: run a notebook using packages in a container

- Working with [Anaconda](#) via an [Apptainer](#) container
  - Build an Apptainer container based from an Anaconda3 Docker container on [Dockerhub](#)
    - Docker pull command: `docker pull continuumio/anaconda3`

```
[user@gra-login]$ module load apptainer  
[user@gra-login]$ apptainer build anaconda3.sif docker://continuumio/anaconda3
```

- Request an interactive session on a compute node via `salloc`
- Load the apptainer module and access the container shell

```
[user@gra-login]$ salloc --account=def-user --cpus-per-task=4 --mem=16G --time=1:0:0  
[user@gra-####]$ module load apptainer  
[user@gra-####]$ apptainer shell anaconda3.sif  
Apptainer>
```

# Use Cases

## Case 3: run a notebook using packages in a container

- Launch a Jupyter server via `jupyter-lab` command from the container shell

```
Apptainer> jupyter-lab --no-browser --ip $(hostname -f)
... ..
http://<node-hostname>:8888/lab?token=109c369980ab346bcd129eb...
```

- Setup a SSH tunnel on a second terminal from your local computer

```
[you@local]$ ssh -L 9999:<node-hostname>:8888 user@cluster.computecanada.ca
[user@cluster-login]$
```

- The URL to connect on a local web browser would be:

<http://localhost:9999/lab?token=109c369980ab346bcd129eb...>